

Cloud composer orchestrating Spark jobs on GKE using the Spark Operator - PoC

[Introduction](#)

[Overview](#)

[Technologies Stack](#)

[Installation of GKE Cluster and Spark Operator](#)

[Deploy a kubernetes Cluster](#)

[Wait for the cluster to be provisioned and running.](#)

[Inspect the kubeconfig file](#)

[Install the Spark Operator](#)

[clone spark-operator repo](#)

[Install the spark operator using the helm chart from the git repo cloned.](#)

[Deploy a test Spark Application](#)

[Generate the connectionUri for authn/Authzn handshake](#)

[Create a kubernetes Secret](#)

[Derive the Client Token](#)

[Build a client token based kubeconfig](#)

[Wrap the Kubeconfig into airflow ready ConnectionUri](#)

[Composer Environment Setup](#)

[Pre-requisites](#)

[Enable the Secret Manager API](#)

[Set up IAM for the PoC](#)

[Setup Secret Manager to host the connection_uri to access the Spark Operator on the GKE cluster](#)

[Create Secret on Secret Manager](#)

[Validate Secret](#)

[Create Composer Environment](#)

[Setup Secret Manager as the backend](#)

[Create Composer Environment](#)

[Validate the Secret Backend](#)

[Familiarise with the DAG and SparkApplications](#)

[Deploy the DAG](#)

[Monitor the Orchestration](#)

Introduction

This page is document the setup involved in getting the GCP cloud composer to orchestrate Spark applications on a remote GKE cluster. The composer uses the SparkKubernetesOperator to perform the job submission to the tenant GKE cluster. For this handshake to happen the composer environment needs to provide a kubernetes secret to the Kube API server/GKE control plane to get the request authenticated and authorised. Such sensitive information needs a dedicated secrets engine and Composer uses the Google Secret Manager as a secret's backend in this setup.

This proof of concept exercise designs the topology in a way that isolates the composer and GKE environments as much as possible to avoid any overlap of networking and security constructs that could be configured as default and stealthily working under the hood to oversimplify the setup. To achieve this the following considerations were made.

- Composer and Spark GKE cluster are isolated over two separate GCP projects (in two different regions)

- An imaginary IaC and Configuration management tool performing the resource provisioning and bootstrapping the secrets across these environments.
- To keep it simple, the traffic between the two projects traverses the internet (however it is encrypted) .

Overview

The use case is where the developer produces Kubernetes ready Spark applications either using the inline coding technique or based on templates and need them orchestrated using a cloud composer environment. This orchestration is documented as Directed Acyclic Graph (DAG) that the composer understands and actions each step by engaging the appropriate operators to perform the task.

SparkKubernetesOperator is a purpose built Composer (Airflow) operator that is capable of encapsulating the mechanics of submitting Spark Applications to a Kubernetes environment. which has Spark Operator up and running and configured to

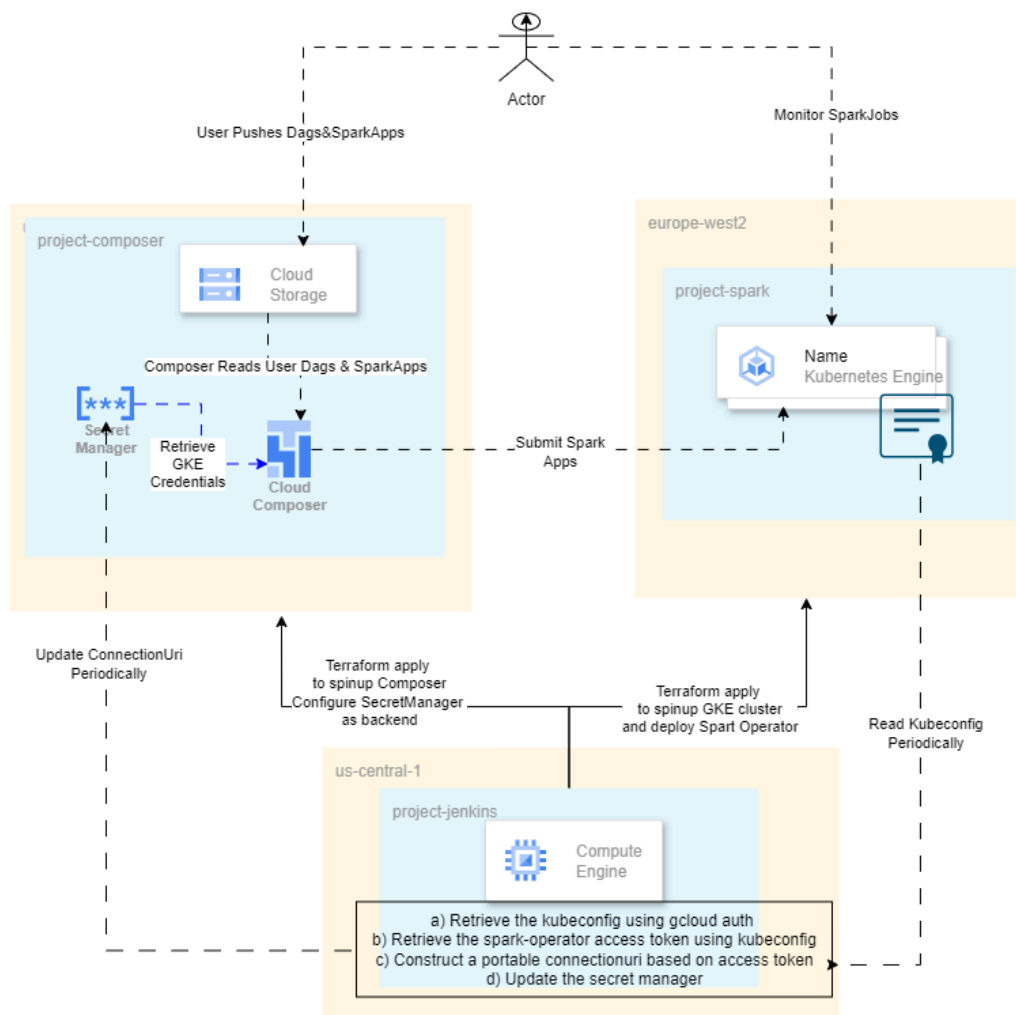
1. Developers write Kubernetes-ready Spark code and DAGs specifying the orchestration requirements for Spark jobs.
2. The Spark code and DAGs are stored in a Google Cloud Storage (GCS) bucket.
3. Cloud Composer, configured to use Google Secret Manager as its secrets backend, reads the DAGs from the GCS bucket.
4. For Spark Kubernetes operators, Cloud Composer packages job requests and fetches secrets from Google Secret Manager to connect to the Kubernetes cluster.
5. The Kubernetes cluster's API Server receives the job request and engages the Spark operator to spin up a Spark application on the Kubernetes cluster.

Additionally there is the infrastructure provisioning aspect where a jenkins like host will take the responsibility of (as shown in the animation below)

- a) Provisioning the Composer
- b) Bootstrapping the Composer with the Secret Manager
- c) Deploying the Tenant Kubernetes Cluster
- d) Deploying the Spark Applications and associated security configuration.
- e) Derive the connectionUri that the composer need to authenticate with GKE and make it available in the Secret manager.

(Note the connectionUri is based on Tokens that are short lived, so there needs to be a recurring process that rotates the secrets and reprovisions the ConnectionUri

Also consideration on the RACE conditions and dependency where the connectionUri creation doesn't impact the SparkJobSubmissions inflight)



Technologies Stack

product	Relevance
Cloud Composer	Google Manager Airflow orchestration tool
Google Secret Manager	A fully managed Secret store and secret engine by Google
Google Kubernetes Engine	Google's Managed Kubernetes Clusters
Spark Operator	Intially a Google curated project to develop an execution platform for Spark jobs on Kubernetes Environment. This is now being managed by the kubeflow community

Installation of GKE Cluster and Spark Operator

```
burnergcp@cloudshell:~ (spark-419510)$ gcloud config list
[accessibility]
screen_reader = True
[component_manager]
disable_update_check = True
[compute]
gce_metadata_read_timeout_sec = 30
[core]
account = burnergcp@gmail.com
disable_usage_reporting = True
project = spark-419510
[metrics]
environment = devshell

Your active configuration is: [cloudshell-31147]
burnergcp@cloudshell:~ (spark-419510)$
```

Deploy a kubernetes Cluster

```
1 export PROJECT_ID=`gcloud config get project`
2
3 #create a gke-cluster
4 gcloud beta container --project $PROJECT_ID clusters create "spark" --machine-type "n2-standard-2" --disk-type "
```

Wait for the cluster to be provisioned and running.

```

gcloud beta container --project=spark-4195101 create a gke-cluster
--node-labels container.googleapis.com/project=spark --machine-type "n2-standard-2" --disk-type "pd-standard" --disk-size "75" --num-nodes "4" --zone "europe-west-2-b" --workload-pool PROJECT_ID.svc.id.goog
--cluster-version "1.21-gke.1500". To create advanced routes based clusters, please pass the "--enable-ip-alias" flag.
Note: Your Rod address range ("--cluster-ip-v4-cidr") can accommodate at most 1008 node(s).
Note: The "europe-west-2-b" Cluster is being health-checked (master is healthy)...done.
Created [https://container.googleapis.com/v1beta1/projects/spark-4195101/zones/europe-west-2/clusters/spark]
To inspect the contents of your cluster, go to: https://console.cloud.google.com/kubernetes/workload/gcloud/europe-west-2-b/spark?project=spark-4195101
kubectl entry generated for spark.

NAME: spark
LOCATION: europe-west-2-b
MASTER VERSION: 1.27.8-gke-1067004
MASTER IP: 34.39.4.214
MACHINE TYPE: n2-standard-2
NODE VERSION: 1.27.8-gke-1067004
NO BOOTDISK: 4
STATUS: RUNNING

```

Cluster create command should get the GKE kubeconfig by default and place it in the default `~/.kube/config`

```
burnergcp@cloudshell:~ (spark-419510) $ ls -l -rt .kube/config
-rw----- 1 burnergcp burnergcp 5525 Apr  8 21:47 .kube/config
burnergcp@cloudshell:~ (spark-419510) $
```

Inspect the kubeconfig file

[illegible]

Pay attention to the user section, there is no x509 crypto like we find in a regular kubernetes kubeconfig file.

The default GKE kubeconfig delegates authorisation to the 'gke-gcloud-auth-plugin' and seeks a jwt token once the oauth2 dance is performed successfully.

This implies the client (either REST APIs, or kubectl in most cases) need to have a means to authenticate with the Google IAM. (we will get back to this aspect later in the demo)

Install the Spark Operator

Now is the time to do some preparation

Create a namespace for the spark-operator

```
1 kubectl create ns spark-operator
```

```
burnergcp@cloudshell:~ (spark-419510) $ #create namespace for spark operator
kubectl create ns spark-operator
namespace/spark-operator created
burnergcp@cloudshell:~ (spark-419510) $ kubectl get ns
NAME                STATUS   AGE
default             Active   22m
gmp-public          Active   21m
gmp-system          Active   21m
kube-node-lease     Active   22m
kube-public         Active   22m
kube-system         Active   22m
spark-operator      Active   14s
burnergcp@cloudshell:~ (spark-419510) $
```

Create a service account and assign it the cluster edit role (which is quite permissive, but we'll keep it simple for the PoC)

```
1 kubectl create sa -n spark-operator spark
```

```
burnergcp@cloudshell:~ (spark-419510) $ kubectl create sa -n spark-operator spark
serviceaccount/spark created
```

Inspect the service account created

```
burnergcp@cloudshell:~ (spark-419510) $ k get sa -n spark-operator spark
NAME      SECRETS   AGE
spark     0         2m25s
burnergcp@cloudshell:~ (spark-419510) $
```

It appears to be created but there are no associated secrets created by default.

(This auto creation of a token has been disabled since the v1.24)

Create a cluster role binding to assign the cluster editor (This is very permissive but is ok for the PoC)

```
1 kubectl create clusterrolebinding spark-role --clusterrole=edit --serviceaccount=spark-operator:spark -n spark-op
```

```
burnergcp@cloudshell:~ (spark-419510) $ kubectl create clusterrolebinding spark-role --clusterrole=edit --serviceaccount=spark-operator:spark -n spark-operator
clusterrolebinding.rbac.authorization.k8s.io/spark-role created
burnergcp@cloudshell:~ (spark-419510) $
```

We now have the kubernetes security resource the service account for the pods to inherit which will empower them with the cluster edit role.

clone spark-operator repo

```
1 git clone https://github.com/kubeflow/spark-operator
```

Note: The spark-operator repository has been moved from GoogleCloud repo and being handled by the kubeflow community

```
burnergcp@cloudshell:~ (spark-419510) $ git clone https://github.com/kubeflow/spark-operator
Cloning into 'spark-operator'...
remote: Enumerating objects: 7487, done.
remote: Counting objects: 100% (352/352), done.
remote: Compressing objects: 100% (232/232), done.
remote: Total 7487 (delta 137), reused 233 (delta 98), pack-reused 7135
Receiving objects: 100% (7487/7487), 24.95 MiB | 24.12 MiB/s, done.
Resolving deltas: 100% (5087/5087), done.
burnergcp@cloudshell:~ (spark-419510) $
```

Install the spark operator using the helm chart from the git repo cloned.

```
1 #Install helm chart for spark operator
2 helm install spark-operator spark-operator/charts/spark-operator-chart --namespace spark-operator
```

This will take a couple of minutes and should install the spark-operator

```
burnergcp@cloudshell:~ (spark-419510) $ #Install helm chart for spark operator
helm install spark-operator spark-operator/charts/spark-operator-chart --namespace spark-operator
NAME: spark-operator
LAST DEPLOYED: Tue Apr 9 07:17:04 2024
NAMESPACE: spark-operator
STATUS: deployed
REVISION: 1
TEST SUITE: None
burnergcp@cloudshell:~ (spark-419510) $
```

Check the objects in the spark-operator namespace

```
burnergcp@cloudshell:~ (spark-419510) $ kubectl get all -n spark-operator
NAME                                READY   STATUS    RESTARTS   AGE
pod/spark-operator-77446ddc6c-2dctn 1/1     Running   0           99s

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/spark-operator      1/1     1            1           99s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/spark-operator-77446ddc6c 1         1         1       99s
burnergcp@cloudshell:~ (spark-419510) $
```

Check the spark custom resource definitions installed as part of the spark operator

```
burnergcp@cloudshell:~ (spark-419510) $ kubectl get crd|grep spark
scheduledsparkapplications.sparkoperator.k8s.io 2024-04-09T07:17:00Z
sparkapplications.sparkoperator.k8s.io         2024-04-09T07:17:00Z
burnergcp@cloudshell:~ (spark-419510) $
```

Now let's run a sample spark application to check if the spark operator works as expected.

The spark-operator repo provides some example applicaitons, however the image uses points to 'gcr.io/*' which is now deprecated

```
burnergcp@cloudshell:~/spark-operator/examples (spark-419510) $ grep 'image:' *
spark-operator-with-metrics.yaml: image: gcr.io/spark-operator/spark-operator:v1beta2-1.3.0-3.1.1
spark-operator-with-webhook.yaml: image: gcr.io/spark-operator/spark-operator:v1beta2-1.3.0-3.1.1
spark-operator-with-webhook.yaml: image: gcr.io/spark-operator/spark-operator:v1beta2-1.3.0-3.1.1
spark-pi-configmap.yaml: image: "gcr.io/spark-operator/spark:v3.1.1"
spark-pi-custom-resource.yaml: image: "gcr.io/spark-operator/spark:v3.1.1"
spark-pi-prometheus.yaml: image: "gcr.io/spark-operator/spark:v3.1.1-gcs-prometheus"
spark-pi-schedule.yaml: image: "gcr.io/spark-operator/spark:v3.1.1"
spark-pi.yaml: image: "gcr.io/spark-operator/spark-py:v3.1.1"
spark-py-pi.yaml: image: "gcr.io/spark-operator/spark-py:v3.1.1"
burnergcp@cloudshell:~/spark-operator/examples (spark-419510) $
burnergcp@cloudshell:~/spark-operator/examples (spark-419510) $
```

The docker pull fails suggested the images are no more available on the gcr registry (These are being ported to the Artifact Registry.)

```
burnergcp@cloudshell:~/spark-operator/examples (spark-419510) $ docker pull gcr.io/spark-operator/spark-py:v3.1.1
Error response from daemon: manifest for gcr.io/spark-operator/spark-py:v3.1.1 not found: manifest unknown: Failed to fetch "v3.1.1" from request "/v2/spark-operator/spark-py/manifests/v3.1.1".
burnergcp@cloudshell:~/spark-operator/examples (spark-419510) $
```

At this point a custom Spark image can be built or use 'fstkfupg/spark-poc-shyam:3.1.1'.

The example we are going to use is 'spark-operator/examples/spark-pi.yaml'. Let's replace the image and the namespace (to point to spark-operator)

```

1 #Update the image and namespace
2 sed -i 's/gcr.io\/spark-operator\/spark:v3.1.1/fstkfupg\/spark-poc-shyam:3.1.1/g' spark-operator/examples/spark-p
3 sed -i 's/namespace\: default/namespace\: spark-operator/g' spark-operator/examples/spark-pi.yaml

```

Now that we have the image changed let's create the spark application.

Deploy a test Spark Application

kubectl apply -f spark-operator/examples/spark-pi.yaml

```

burnergcp@cloudshell:~ (spark-419510) $ k get sparkapplications.sparkoperator.k8s.io -n spark-operator
NAME          STATUS    ATTEMPTS  START              FINISH              AGE
spark-pi      COMPLETED 1          2024-04-09T07:20:02Z 2024-04-09T07:20:58Z 5m44s
burnergcp@cloudshell:~ (spark-419510) $

```

The Spark Application will launch a Driver and an Executor tasks and finally runs to completion,

Note the sequence of events (SparkApplicationAdded → SparkApplicationSubmitted → SparkDriverRunning → SparkExecutorPending → SparkExecutorRunning → SparkDriverCompleted → SparkApplicationCompleted → SparkExecutorCompleted)

```

kubectl describe sparkapplications -n spark-operator |grep 'Events:' -A10
sparkapplication.sparkoperator.k8s.io/spark-pi created
NAME          STATUS    ATTEMPTS  START              FINISH              AGE
spark-pi      COMPLETED 1          2024-04-09T07:20:02Z 2024-04-09T07:20:58Z 0s
Events:
Type      Reason                                     Age    From          Message
-----
Normal    SparkApplicationAdded                    2ms    spark-operator SparkApplication spark-pi was added, enqueueing it for submission
Normal    SparkApplicationSubmitted                116s   spark-operator SparkApplication spark-pi was submitted successfully
Normal    SparkDriverRunning                      95s    spark-operator Driver spark-pi-driver is running
Normal    SparkExecutorPending                    87s    spark-operator Executor [spark-pi-14bcf88ec1bacd1d-exec-1] is pending
Normal    SparkExecutorRunning                    66s    spark-operator Executor [spark-pi-14bcf88ec1bacd1d-exec-1] is running
Normal    SparkDriverCompleted                    60s    spark-operator Driver spark-pi-driver completed
Normal    SparkApplicationCompleted                60s    spark-operator SparkApplication spark-pi completed
Normal    SparkExecutorCompleted                   58s    spark-operator Executor [spark-pi-14bcf88ec1bacd1d-exec-1] completed
burnergcp@cloudshell:~ (spark-419510) $

```

Generate the connectionUri for authn/Authzn handshake

So far we have been running these using kubectl which in turn uses the .kube/config that is provided as part of the cluster creation. While the a regular kubeconfig file has enough details, i.e (cluster details, CA certs, Users details and User Tokens) to authenticate with a kubeAPI server the version that's provided by the GKE cluster federates the user auth with gloud-auth plugin. So this Bring Your Own Auth sort of kubeconfig will not suffice to authenticate from a client that can't integrate with gloud-auth-plugin. For this reason we need to find a way to get the composer operator authenticate with the GKE control plane over the client certificates.

This can be achieved by creating a dedicated secret 'airflow-secret' and bootstrapping with the 'spark-operator' Kubernetes Service Account.

Create a kubernetes Secret

```

1 apiVersion: v1
2 kind: Secret
3 metadata:
4   name: airflow-secret
5   namespace: spark-operator
6   annotations:
7     kubernetes.io/service-account.name: spark-operator
8 type: kubernetes.io/service-account-token
9
10

```



```
burnergcp@cloudshell:~ (spark-419510) $ kubectl config view -ojson --raw > kubeconfig
burnergcp@cloudshell:~ (spark-419510) $
burnergcp@cloudshell:~ (spark-419510) $ ls -lrt kubeconfig
-rw-r--r-- 1 burnergcp burnergcp 3741 Apr  8 23:33 kubeconfig
burnergcp@cloudshell:~ (spark-419510) $
```

We need the use the airflow connection library to derive the connection_uri. thankfully there is a straight forward function `Connection.get_uri()` we can leverage.

Following is a quick script to print the connection_uri as expected by the composer. This connection_uri needs to made available to the composer environment , hence placed in the secret manager so that the airflow DAG can retrieve it when needed.

Wrap the Kubeconfig into airflow ready ConnectionUri

```
1 import json
2 from airflow.models.connection import Connection
3
4 with open('/home/burnergcp/kubeconfig') as f:
5     kube_config_json = json.load(f)
6
7 c = Connection(
8     conn_id='sparkgke',
9     conn_type='kubernetes',
10    extra=json.dumps(dict(
11        in_cluster="false",
12        kube_config=kube_config_json,
13        namespace='spark-operator'))
14    )
15 print(c.get_uri())
```

```
burnergcp@cloudshell:~ (spark-419510) $ python getconn_uri.py
1 #!/usr/bin/env python
2 """
3 This script will generate a kubeconfig file for the spark operator.
4 """
5 import sys
6 import os
7 import json
8 import base64
9 import subprocess
10
11 # Get the kubeconfig file
12 kubeconfig_file = '/home/burnergcp/kubeconfig'
13
14 # Read the kubeconfig file
15 with open(kubeconfig_file) as f:
16     kubeconfig_json = json.load(f)
17
18 # Create a Connection object
19 conn_id = 'sparkgke'
20 conn_type = 'kubernetes'
21 extra = {
22     'in_cluster': 'false',
23     'kube_config': kubeconfig_json,
24     'namespace': 'spark-operator'
25 }
26
27 # Create the Connection
28 conn = Connection(conn_id=conn_id, conn_type=conn_type, extra=json.dumps(extra))
29
30 # Print the URI
31 print(conn.get_uri())
```

That is all we need to do on the GKE side. We now move on to the Composer project to setup the Composer environment

Composer Environment Setup

Pre-requisites

Enable the Secret Manager API

so that we can use the product in the first place

```
burnergcp@cloudshell:~ (composer-13322) $ gcloud services enable secretmanager.googleapis.com
Operation "operations/acet.p2-339838880809-e2c26e41-5fd4-45de-89d5-0a530c398bb9" finished successfully.
burnergcp@cloudshell:~ (composer-13322) $
```

The composer environment needs access to fetch secrets from Secret Manager, so it needs the 'Secret Accessor Role'. As we are going to create the composer environment with the default compute SA in the PoC, granting the role the same.

Never do this production, always use a dedicated SA with curated permissions

Set up IAM for the PoC

New principals *

339838880809-compute@developer.gserviceaccount.com ✕ ?

Assign roles

Roles are composed of sets of permissions and determine what the principal can do with this resource. [Learn more](#)

Role *
Secret Manager Secret Accessor ▼

IAM condition (optional) ?
[+ ADD IAM CONDITION](#)

Allows accessing the payload of secrets.

[+ ADD ANOTHER ROLE](#)

[SAVE](#) [CANCEL](#)

Setup Secret Manager to host the connection_uri to access the Spark Operator on the GKE cluster

The connection_uri we managed to create using the client token based kubeconfig needs to be persisted in the Secret Manager. Composer is a bit touchy about the naming convention when it comes to fetching the actual connection. It expects the secret name to be prefixed with 'airflow-connections-'

Create a new secret 'airflow-connections-sparkgke'. where 'airflow-connections-' is the default prefix sparkgke is our userdefined friendly connection_id name

The prefix configuration can be altered using some overrides but who cares for this PoC, we will stick to the defaults and use 'airflow-connections-'

So the composer during the execution time will expect the 'kubernetes_conn_id' as 'sparkgke' prepends 'airflow-connections-' and tries to resolve the secret and fetches it to kick off the spark operator.

et

with the secret value in the first version

parkgke

identifiable and unique within this project.

or import it directly from a file.

?

!-b_spark%22%7D%7D%5D%2C+%22c
gke_spark-419510_europe-west2-
amespace=spark-operator

FE10D0D6 ?

y

ocations for this secret

Encryption key
paired

CANCEL

Use the following command to confirm the secret is created as expected with no typos, spaces or tabs

[illegible]

Create a cloud composer environment

Leave the service account default to the compute developer (as discussed earlier), however it needs Cloud Composer V2 API Service Agent Extension role as it complains below. This can be set from the same screen without having to navigate to IAM screen or using gcloud command.

Check the box and click on Grant while the warning disappears

Name *

highcpu

Location *


us-east4

Image version *

composer-2.6.6-airflow-2.6.3

Service account

1086564652743-compute@developer.gserviceaccount.com



Grant required permissions to Cloud Composer service account

Cloud Composer relies on [Workload Identity](#) as Google API authentication mechanism for Airflow.

In order to support Workload Identity, Cloud Composer creates additional IAM role bindings which requires the [Cloud Composer v2 API Service Agent Extension](#) role.

☒ Grant the Cloud Composer v2 API Service Agent Extension role to the service-1086564652743@cloudcomposer-accounts.iam.gserviceaccount.com service account.
Service account 1086564652743-compute@developer.gserviceaccount.com will be used as a resource.

GRANT

Labels ?

Now that the Cloud Composer V2 API Service Agent Extension role is assigned to the service account, composer create screen IAM validation is now OK.

Name *

highcpu

Location *

us-east4

Image version *

composer-2.6.6-airflow-2.6.3

Service account

1086564652743-compute@developer.gserviceaccount.com

Labels ?

+ ADD LABEL

Resilience mode

- ☒ Standard resilience (default)
- ☐ High resilience

Setup Secret Manager as the backend

The Composer needs an additional setting in the 'Advanced Configuration' section. This is to configure the backend where composer stores and retrieves its sensitive information. As the aim of the PoC is to prove that composer can use Google Secret Manager as the secret backend, this needs to be set up accordingly.

Network configuration

The network configuration for the Google Kubernetes Engine cluster running the Airflow software.

SHOW NETWORK CONFIGURATION

Advanced configuration

Environment variables, Airflow configuration overrides, encryption, and maintenance.

SHOW ADVANCED CONFIGURATION

In the 'Airflow configuration overrides' add a new configuration override

- Section 1: secrets
- Key 1: backend
- Value 1: airflow.providers.google.cloud.secrets.secret_manager.CloudSecretManagerBackend

Airflow configuration overrides ?

Section 1 *
secrets

Key 1 *
backend

Value 1 *
cretManagerBackend

+ ADD AIRFLOW CONFIGURATION OVERRIDE

Click the 'Create' button to create the composer environment

Create Composer Environment

The composer starts to spin up and might take 15-20 mins.

Composer

Environments

CREATE

REFRESH

DELETE

Filter

Filter environments

State	Name	Location	Composer version	Airflow version	Creation time	Update time	Airflow webserver	DAG list	Logs	DAGs folder	Labels
	highcpu	us-central1	2.6.6	2.6.3	4/9/24, 10:10 AM	4/9/24, 10:10 AM	None	DAGs	Logs	None	None

Composer

Environments

CREATE

REFRESH

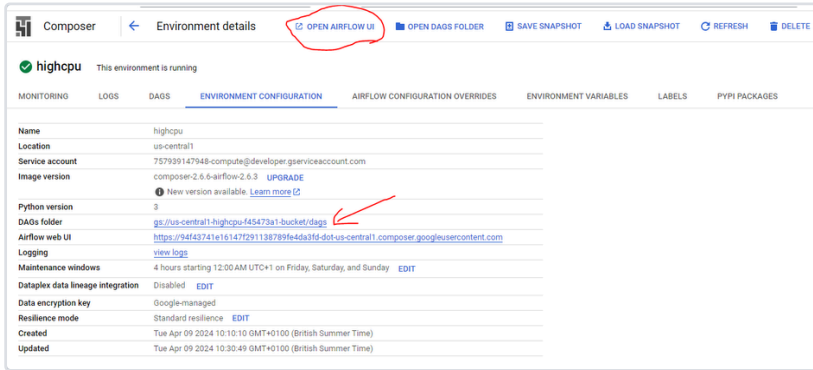
DELETE

Filter

Filter environments

State	Name	Location	Composer version	Airflow version	Creation time	Update time	Airflow webserver	DAG list	Logs	DAGs folder	Labels
	highcpu	us-central1	2.6.6	2.6.3	4/9/24, 10:10 AM	4/9/24, 10:30 AM	Airflow	DAGs	Logs	DAGs	None

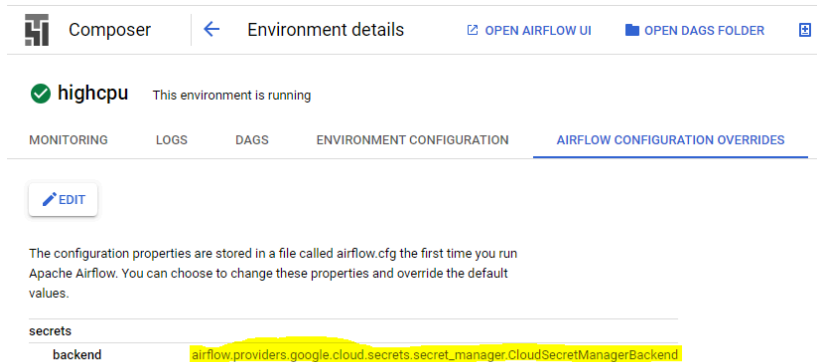
Once the Composer environment is ready, it provides the configuration as to where it fed the DAGs and application code from, Its the GCS bucket as listed below. Also the AIRFLOW UI button get enabled.



Validate the Secret Backend

Before you go ahead, it's important to ensure the secrets backed for the composer is indeed the Google Secrets Manager, this can be checked from the UI or using CLI

On the UI look for the 'AIRFLOW CONFIGURTAION OVERRIDES'



Or use the following gcloud command to list the backend

```
1 gcloud composer environments describe highcpu --location=us-east1|grep secrets
```

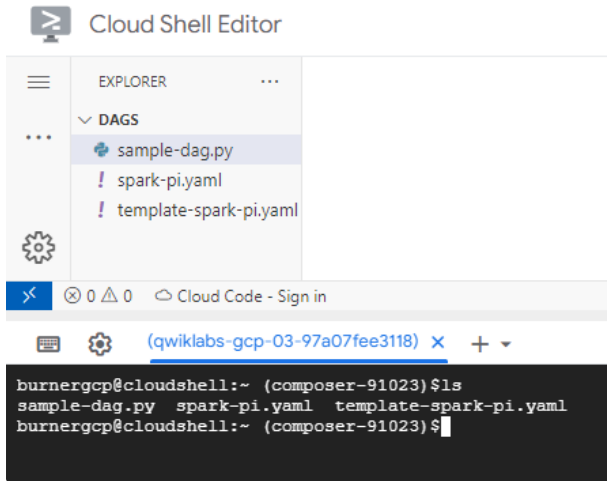
```
burnergop@cloudshell:~ (composer-91023) $  
burnergop@cloudshell:~ (composer-91023) $ gcloud composer environments describe highcpu --location=us-east1|grep secrets  
secrets: backend: airflow.providers.google.cloud.secrets.secret_manager.CloudSecretManagerBackend  
burnergop@cloudshell:~ (composer-91023) $
```

Once we are happy with the backend settings, let's move on to deploying the DAGs along with the associated application code.

Familiarise with the DAG and SparkApplications

For this PoC, let's use the same spark-py.yaml based Spark Application we deployed using the Kubectl locally on the tenant GKE cluster. We need a DAG to orchestrate the SparkOperator to be able to submit the job to the Cluster

As the PoC's aim is to highlight the Integration between the Composer and Spark Operator we chose the simplest scala application for test.



File	Purpose	Content
sample-dag.py	DAG to orchestrate the workflow	<div>Click here to expand...</div> <pre>1 """ 2 This is an example 3 In this example, w 4 The first task is 5 and the second tas 6 7 Spark-on-k8s opera 8 https://github.com 9 """ 10 11 from datetime impo 12 13 # [START import_mo 14 # The DAG object; 15 from airflow impor 16 # Operators; we ne 17 from airflow.opera 18 from airflow.provi 19 from airflow.provi 20 from airflow.provi 21 from airflow.utils 22 #k8s_hook = Kuberr 23 # [END import_modu 24 25 # [START default_e 26 # These args will 27 # You can override 28 default_args = { 29 'owner': 'gkes 30 'depends_on_pa 31 'start_date': 32 } 33 # [END default_arg 34 35 # [START instantia 36 37 dag = DAG(38 'spark_demo', 39 start_date=day</pre>

		<pre>40 default_args=c 41 schedule_inter 42 tags=['secmana 43) 44 # Define the BashC 45 hello_world_task = 46 task_id='hello 47 bash_command=' 48 dag=dag 49) 50 51 # Define the task 52 53 submit = SparkKube 54 task_id='spark 55 namespace='spa 56 application_fi 57 kubernetes_cor 58 do_xcom_push=T 59) 60 61 62 hello_world_task > 63</pre>
template-spark-pi.yaml	A template file, basically a pointer to the actual Spark Application manifest	<pre>1 /home/airflow/gcs/dags/sp</pre>
spark-pi.yaml	The yaml that contains the Spark Application manifest as provided in the Spark Operator Examples	<div>▼ Click here to expand...</div> <pre>1 apiVersion: "spark 2 kind: SparkApplica 3 metadata: 4 name: spark-pi 5 namespace: spark 6 spec: 7 type: Scala 8 mode: cluster 9 # Replce with yo 10 image: "fstkfupc 11 imagePullPolicy: 12 mainClass: org.æ 13 mainApplicationF 14 sparkVersion: "3 15 restartPolicy: 16 type: Never 17 volumes: 18 - name: "test- 19 hostPath: 20 path: "/tn 21 type: Dire 22 driver: 23 cores: 1 24 coreLimit: "12 25 memory: "512m" 26 labels: 27 version: 3.1 28 serviceAccount</pre>

```

29     volumeMounts:
30       - name: "tes
31         mountPath:
32     executor:
33       cores: 1
34       instances: 1
35       memory: "512m"
36     labels:
37       version: 3.1
38     volumeMounts:
39       - name: "tes
40         mountPath:
41

```

Deploy the DAG

Copy the dag and associated spark code to the gcs bucket that the composer is configured to

```

burnergcp@cloudshell:~ (composer-91023)$ ls -lrt
total 4
drwxr-xr-x 2 student_01_01625eba5c88 student_01_01625eba5c88 4096 Apr  9 08:37 dags
burnergcp@cloudshell:~ (composer-91023)$ cd dags
burnergcp@cloudshell:~ (composer-91023)$ ls -lrt
total 12
-rw-r--r-- 1 student_01_01625eba5c88 student_01_01625eba5c88 37 Apr  9 08:36 template-spark-pi.yaml
-rw-r--r-- 1 student_01_01625eba5c88 student_01_01625eba5c88 914 Apr  9 08:36 spark-pi.yaml
-rw-r--r-- 1 student_01_01625eba5c88 student_01_01625eba5c88 2027 Apr  9 08:37 sample-dag.py
burnergcp@cloudshell:~ (composer-91023)$
burnergcp@cloudshell:~ (composer-91023)$ gsutil cp * gs://us-central1-highcpu-f45473a1-bucket/dags
Copying file:///sample-dag.py [Content-Type=text/x-python]...
Copying file:///spark-pi.yaml [Content-Type=application/octet-stream]...
Copying file:///template-spark-pi.yaml [Content-Type=application/octet-stream]...
\ [3 files] [ 2.9 KiB/ 2.9 KiB]
Operation completed over 3 objects/2.9 KiB.
burnergcp@cloudshell:~ (composer-91023)$

```

They seem to be available in the dags folder on the GCS bucket Composer reads from

us-central1-highcpu-f45473a1-bucket Updated just now

Location	Storage class	Public access	Protection
us-central1 (Iowa)	Standard	Subject to object ACLs	None

OBJECTS CONFIGURATION PERMISSIONS PROTECTION LIFECYCLE OBSERVABILITY INVENTORY REPORTS OPERATIONS

Folder browser

Buckets > us-central1-highcpu-f45473a1-bucket > dags

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER TRANSFER DATA MANAGE HOLDS EDIT RETENTION DOWNLOAD

DELETE


Filter by name prefix only Filter Filter objects and folders Show Live objects only


<input type="checkbox"/>	Name	Size	Type	Created	Storage class	Last modified	
<input type="checkbox"/>	airflow_monitoring.py	809 B	text/x-python	Apr 9, 2024, 10:30:08 AM	Standard	Apr 9, 2024, 11	
<input type="checkbox"/>	sample-dag.py	2 KB	text/x-python	Apr 9, 2024, 10:31:12 AM	Standard	Apr 9, 2024, 11	
<input type="checkbox"/>	spark-pi.yaml	914 B	application/octet-stream	Apr 9, 2024, 10:31:13 AM	Standard	Apr 9, 2024, 11	
<input type="checkbox"/>	template-spark-pi.yaml	37 B	application/octet-stream	Apr 9, 2024, 10:31:14 AM	Standard	Apr 9, 2024, 11	

The worker nodes of composer will eventually read the DAGs and expose them on the UI (note that the tags and owner metadata is reflected on the UI for ease of management)










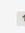
Admin ▾ Docs ▾ 



 Airflow DAGs Datasets Browse Admin Docs Composer


highcpu

All 2
Active 2
Paused 0

DAG	Owner	Runs	Schedule	Last
airflow_monitoring	airflow	   	*/*/* * * * *	202
spark_demo	airflow	   	1 day, 0:00:00	202

« < 1 > »

Submit a manual run



Clean up the Spark applications in the spark-operator namespace

```

burnergcp@cloudshell:~ (spark-419510) $ k get sparkapplications.sparkoperator.k8s.io -n spark-operator
No resources found in spark-operator namespace.
burnergcp@cloudshell:~ (spark-419510) $
  
```

We now see the Composer DAG's SparkKubernetesOperator submitting the spark-py.yaml using the connection_id 'sparkgke'

```

burnergcp@cloudshell:~ (spark-419510) $ k get sparkapplications.sparkoperator.k8s.io -n spark-operator
NAME                                STATUS    ATTEMPTS   START                FINISH              AGE
spark-job-composer-xt9riotq         RUNNING   1          2024-04-09T13:32:43Z <no value>         21s
burnergcp@cloudshell:~ (spark-419510) $
  
```

```
burnergcp@cloudshell:~ (spark-419510) $ k get pods -n spark-operator
```

NAME	READY	STATUS	RESTARTS	AGE
spark-job-composer-gw5gnszc-driver	0/1	Completed	0	19s
spark-job-composer-xt9riotq-driver	0/1	Completed	0	107s
spark-operator-77446ddc6c-ctwdn	1/1	Running	2 (92m ago)	93m

```
burnergcp@cloudshell:~ (spark-419510) $
```